Formal Verification of Session Types

by

Théa Johnson,

Thesis

Presented to the

University of Dublin, Trinity College

in fulfillment

of the requirements

for the Degree of

Bachelor's of Arts

University of Dublin, Trinity College

April 2016

Declaration

I, the undersigned, declare that this work has not previously been submitted as an exercise for a degree at this, or any other University, and that unless otherwise stated, is my own work.

Théa Johnson

April 18, 2016

Permission to Lend and/or Copy

I, the undersigned, agree that Trinity College Library may lend or copy this thesis upon request.

Théa Johnson

April 18, 2016

Formal Verification of Session Types

Théa Johnson, B.A.
University of Dublin, Trinity College, 2016

Supervisor: Dr. Vasileios Koutavas

...ABSTRACT...

Contents

Chapter 1

Introduction

1.1 Project Objectives

The aims of this project are to investigate the formal verification of session types, specifically in relation to the paper Type-Based Analysis for Session Inference [1]. Also to implement a behaviour checker based on the designs described in this paper using OCaml, Menhir and OCamllex.

1.2 Summary of Report

In this report first, in this section, a brief background to the area is given. The following section then gives an overview and explanation of the system proposed in the paper Type-Based Analysis for Session Inference [1]. The third section details the implementation completed over the course of this project. The final chapters cover some detailed examples of how programs are dealt with in this system and an evaluation of the system.

1.3 Motivation

The modern world is growing increasingly dependent on distributed systems, changing the historical approach to computing dramatically. In order for modern society to function it is important that these systems communicate correctly and that when proposing

or introducing new systems we can show that they will communicate correctly under all circumstances.

Modern programming languages support data types. These allow us to use verification techniques to show that the program will run as expected on all forms of input. A similar system of types could be used for communication over distributed systems. Ideally a type system for communication would be embedded into languages in a similar fashion to the data type systems of modern languages.

This is the system that is proposed with session types. These types can specify the style of communication expected (in general terms send or receive) as well as the type of data that is expected.

1.4 Current Work

This area is currently been researched by multiple groups. However it is currently not used in real world systems to any great extent. To date systems have been developed for applying session type disciplines to functional languages, object oriented languages and operating systems.

1.5 Background

Traditional type systems embedded into programming languages focus on the computations and what they should produces. Session type disciplines aim for embedded session types that can describe the sequence of messages as well the type of the messages transmitted on communication channels. Then, since the session types will describe the protocol of a channel, verification techniques can be used to ensure that processes will abide by these protocols.

1.5.1 Main session type approaches

The main approaches to session types, according to (Hüttle et al.) [2] are detailed in the following section.

Session types are usually associated with binary communication channels where the two ends using the channel view the endpoints as complementary types. Static type

checking can then be used to ensure that the communications on the channel abide by the protocol specified.

Multiparty session types extend binary session types to allow for more than two processes to communicate.

Conversation session types unify local and global multipary types and allow for an unspecified number of processes to communicate over the channel.

Contracts focus on general theory to confirm that communications follow the specified abstract description of input/output actions.

Chapter 2

Summary of the Paper "Type Based Analysis for Session Inference"

2.1 Overview

This paper[1] proposes a system for a design approach to Binary Session Types which uses effects. For this high level language is developed where communication protocols can be programmed and statically checked. In the paper the approach is applied to ML_s, a core of the language ML with session communication.

The approach suggested separates traditional typing and session typing using a two level system. The first level uses a type and effect system, which is an adaptation of the one developed by (Amtoft, Neilsen and Neilsen, 1999)[3]. The types allow for a clear representation of program properties and enable modular reasoning. The effects are essentially behaviour types that record the actions of interest that will occur at run time. From this system a complete behaviour inference algorithm is obtained, this extracts a behaviour for a program providing it respects ML types. In this level the programs are typed against both an ML type and a behaviour. Session protocols are not considered here and so endpoints (see 2.3) are given a type ses^{ρ} instead.

The second level checks the behaviour to see that it complies with the session types of both channels and endpoints. In performing this check the operational semantics

are used (see 2.3).

This level is inspired by the work done by Castagna et al. [4]. In their system session based interaction is established on a public channel, once established the parties continue to communicate on a private channel. Messages are exchanged on the private channel according to a given protocol. Internal and external choices are also required to implement control. Internal choices are when the decision is made autonomously by a process and external choices occur when a decision is based entirely on messages received.

This level ensures that sessions respect the order of communication and message types described by the session type of the channel. It also ensures partial lock freedom due to stacked interleaving of sessions.

One of the most appealing aspects of the session type discipline proposed here is that it allows for complete session type inference from behaviours. When this is combined with behaviour inference from level 1 a method for complete session type inference without programmer annotations is achieved.

The two levels of the system only interact through behaviours. This allows for the development of front ends for different languages and back ends for different session disciplines and to combine the two to cover an extensive selection of requirements. The behaviour checker implemented (see 3.4) can be used with any implementation of the first level provided that the output is of the correct format.

2.2 The First Level

At this level the type and effect system of (Amtoft, Neilsen and Neilsen, 1999) [3] is extended to session communications in ML_s. The type and effect system consists of constructions of judgments. Judgments are of the form $C; \Gamma \vdash e : T \triangleright b$. In these statements C represents the constraint environment which is used to relate type level variables to terms and so enables session inference. Γ represents the type environment which is used to bind program variables to type schemas. To read this judgment we would say that expression e has type T and behaviour b under type environment Γ and constraint environment C.

In the system designed in the paper an ML_s expression can have either a standard ML type or a session type. Session types are of the form ses^{ρ} where ρ is a static

to fill in

Figure 2.1: Syntax of types, behaviours and constraints

approximation of the location of the endpoint. Functional types have an associated behaviour β and type variables α are used for ML polymorphism.

Polymorphism is extended with type schemas. These are of the form $\forall(\overrightarrow{\gamma} : C_0).T$ where γ is a list made up of some combination of type (α), behaviour (β), region (ρ) and session (ψ) variables and C_0 represents the constraint environment that imposes constraints on the quantified variables.

2.2.1 Syntax of types, behaviours and constraints

Of the types detailed in 2.1 behaviours and constraints are the most relevant to the implementation discussed in 3. Behaviours can be of the simple forms β a behaviour variable, τ the behaviour with no effect, $b; b$ a sequence of behaviours or $b \oplus b$ a choice in behaviours. The more complicate forms also include $rec_\beta b$ for recursive behaviour. In this case, and in the case of *spawnb*, the body of the recursion must be confined (must not effect open endpoints and must consume all endpoints it opens). Behaviour can also input or output types ($\rho!T, \rho?T$), delegate and resume endpoints ($\rho!\rho, \rho!l$), select from an internal choice $\rho!L_i$ and offer external choice $\underset{i \in I}{\&}\{\rho?L_i; b_i\}$.

Constraints can specify that types are subtypes of another type $T \subseteq T$, or that they are confined $cfd(T)$. They also specify what behaviours behaviour variables can act as $b \subseteq \beta$. Region constraints ($\rho \sim r$) link region variables to other region variables or to locations. Channel constraints ($c \sim \eta, \bar{c} \sim \eta$) specify the link between channels and their endpoints. The duality constraint ($\eta \bowtie \eta$) states that two endpoints are complementary. Sets of constraints can also be specified.

Type schemas, locations and region variables

Through the constraint environment (C) region constraints specify links between region variables (ρ) and other region variables or labels. These region constraints are produced during pre-processing. They identify the textual source of endpoints. The labels l specify locations in the input code that endpoints are generated.

For example if $req - c^l$ is called l is the location in the code it is called from. This

$$req - c^l : \forall(\beta\rho\psi : push(l : \psi) \subseteq \beta, \rho \sim l, c \sim \psi).Unit \xrightarrow{\beta} Ses^\rho$$

$$acc - c^l : \forall(\beta\rho\psi : push(l : \psi) \subseteq \beta, \rho \sim l, \bar{c} \sim \psi).Unit \xrightarrow{\beta} Ses^\rho$$

$$ressume - c^l : \forall(\beta\rho\rho' : \rho?\rho' \subseteq \beta, \rho' \sim l).Ses^\rho \xrightarrow{\beta} Ses^{\rho'}$$

$$recv : \forall(\alpha\beta\rho : \rho?\alpha \subseteq \beta cfd(\alpha)).Ses^\rho \xrightarrow{\beta} \alpha$$

$$send : \forall(\alpha\beta\rho : \rho!\alpha \subseteq \beta cfd(\alpha)).Ses^\rho \times \alpha \xrightarrow{\beta} Unit$$

$$deleg : \forall(\alpha\rho\rho' : \rho!\rho' \subseteq \beta).Ses^\rho \times Ses^{\rho'} \xrightarrow{\beta} Unit$$

$$sel - L : \forall(\alpha\rho : \rho?L \subseteq \beta).Ses^\rho \xrightarrow{\beta} Unit$$

Figure 2.2: Type Schemas

means that if this were to be called in, for example, a for loop we would have multiple instances of endpoints related to l. This is a limitation of this system.

If we have an expression with type Ses^ρ which evaluates to p^l then a region constraint must exist to link l to ρ. This takes the form of $C \vdash \rho \sim l$ which says that ρ and l are linked under constraint environment C. This tells us that p was generated from the location int the code referenced by l.

If we were to look up this location we would find one of $req-c^l, acc-c^l or resume-c^l$ where c references the private channel on which the communication will take place. These primitive functions are typed by the rule TConst given in fig. 2.3

They type schemas of these primitives are given in fig. 2.2. In the case of $req - c^l$ a new session is started on the static endpoint l. In order for it to be type-able C must contain its effect which is $push(l : \psi) \subseteq \beta$. In the stack frame ψ is the session variable that represents the session type of endpoint l. C must also record that the region variable ρ is linked to l and that the 'request' endpoint of c, the channel, has session type ψ.

The remaining type schemas are read in a similar way.

2.2.2 Type and effect system

The rules for the type and effect system proposed are given in fig. 2.3. These consist of the judgments described above and requirements for the rule. These rules say that

To fill in

Figure 2.3: Type and Effect system for ML_s Expressions omitting rule for pairs

if we have a judgments of the form given above the line and if the requirements beside the rule (if they exist) are met then the judgment below the line will be valid.

For example the rule TSUB it states that if we have constraint environment C, type environment Γ, expression e of type T with associated behaviour b and we also know that constraint environment C contains constraints telling us that T is a functional subtype of T' and the b is a sub-behaviour of β then we can also say that under the same type and constraint environments the expression e can be said to have type T' and behaviour β.

Of the remaining rules TLeft, TVar, TIf, TConst, TApp, TFun, TSpawn and the rule for pairs are used to perform standard type checking of sequential and non-deterministic behaviour. TSub allows us to replace a behaviour variable with its associated behaviour. TIns and TGen are used to extend the instantiation and generalisation rules of ML.

TRec ensures that in the case of recursion the communication effect of the body of the recursion is confined. This means that it will not effect any endpoints already open when called and will consume all end points opened during its execution.

TEndP ensure that if we have an expression with the type of a session endpoint with associated region and this expression evaluates to a value associated with a location then there exists a link between the region and the location in the constraint environment.

TConst types primitive functions such as $req - c^l$.

2.3 The Second Level

At this level session types are considered. Theses take the form of:

$$\eta ::= end | !T.\eta | ?T.\eta | !\eta.\eta | ?\eta.\eta | \bigoplus_{i \in I} \{L_i : \eta_i\} | \underset{i \in (I_1, I_2)}{\&} \{L_i : \eta_i\} | \psi$$

Either communication on the endpoint is finished (end) or more communications are going to take place. These include sending or receiving a confined type ($!T.\eta, ?T.\eta$),

END:	$(l : end).\Delta \models b \rightarrow c\Delta \models b$	
BETA:	$\Delta \models \beta \rightarrow c\Delta \models b$	if $C \vdash b \subseteq \beta$
PLUS:	$\Delta \models b_1 \oplus b_2 \rightarrow c\Delta \models b_i$	if $i \in 1, 2$
PUSH:	$\Delta \models push(l : \eta) \rightarrow c(l : \eta).\Delta \models \tau$	if $l \notin \Delta.labels$
OUT:	$(l :!T.\eta).\Delta \models \rho!T' \rightarrow c(l : \eta).\Delta \models \tau$	if $C \vdash \rho \sim l, T' <: T$
IN:	$(l :?T.\eta).\Delta \models \rho?T' \rightarrow c(l : \eta).\Delta \models \tau$	if $C \vdash \rho \sim l, T <: T'$
DEL:	$(l :!\eta_d.\eta).(l_d : \eta'_d).\Delta \models \rho!\rho_d \rightarrow c(l : \eta).\Delta \models \tau$	if $C \vdash \rho \sim l, \rho_d \sim l_d, \eta'_d <: \eta_d$
RES:	$(l :?\eta_r.\eta) \models \rho?l_r \rightarrow c(l : \eta).(l_r : \eta_r) \models \tau$	if $(l \neq l_r), C \vdash \rho \sim l$
ICH:	$(l : \underset{i \in I}{\oplus}\{L_i : \eta_i\}).\Delta \models \rho!L_j \rightarrow c(l : \eta_j).\Delta \models \tau$	if $(J \in I), C \vdash \rho \sim l$
ECH:	$(l : \underset{i \in (I_1, I_2)}{\&}\{L_i : \eta_i).\Delta \models \underset{j \in J}{\&}\{\rho?L_j : b_j\} \rightarrow c(l : \eta_k).\Delta \models b_k$	if $k \in J, C \vdash \rho \sim l, I_1 \subseteq J \subseteq I_1 \cup I_2$
REC:	$\Delta \models rec_\beta b \rightarrow c\Delta \models \tau$	if $\epsilon \models b \Downarrow c', C' = (C \backslash (rec_\beta b \subseteq \beta)) \cup (\tau \subseteq \beta)$
SPN:	$\Delta \models spawn b \rightarrow c\Delta \models \tau$	if $\epsilon \models b \Downarrow c$
SEQ:	$\Delta \models b_1; b_2 \rightarrow c\Delta' \models b'_1; b_2$	if $\Delta \models b_1 \rightarrow c\Delta' \models b'_1$
TAU:	$\Delta \models \tau; b \rightarrow c\Delta \models b$	

Figure 2.4: Abstract Interpretation Semantics

delegating by sending one endpoint over another ($!\eta.\eta$) and resuming an endpoint ($?\eta.\eta$). Non deterministic selection involves selecting a label L_i, selected by the selection behaviour, and then following the selected session type η_i. External choice ($\underset{i \in (I_1, I_2)}{\&}\{L_i : \eta_i\}$) is also supported.

2.3.1 Abstract interpretation semantics

In these semantics (fig. 2.4) Δ represents the stack which consists of frames of a label and a session type, c represents the channel, C is the set of constraints. A behaviour and the current stack are taken and an attempt is made to match them to one of the rules. These rules then describe what actions are to be taken.

For example the REC Rule states that if the current behaviour is recursion then we must check that the behaviour associated with that recursion will follow these rules, starting with an empty stack, to end in a state with the empty stack and the behaviour τ. Another constraint is that any behaviour constraints in the constraint environment of the form $rec_\beta b \subseteq \beta$ must be replaced with $\tau \subseteq \beta$. If these constraints are met then

the next behaviour to be checked is τ with the stack remaining unchanged.

This ensures that recursive behaviour is confined, which is necessary for this system.

Of the other rules END removes a finished stack frame. BETA looks up behaviour variables and subs in the behaviours associated with them in the constraint environment. PLUS choses a branch. PUSH adds a new frame to the stack given that the label has not previously been pushed to the stack. OUT and IN reduce the frame at the top of the stack. DEL and RES are used to transfer endpoints. RES must be applied to a one frame stack to ensure that there are no two endpoints of the same session pushed to the same stack (this is to avoid deadlock). ICH offers internal choice of session types. ECH offers external choice. SPN ensures that spawned processes are confined.

2.4 Inference Algorithm

There are three inference algorithms used in the proposed system. The first of these is used with the first level to infer functional types and communication effects. The remaining two are used with the second level to infer session types from the abstract interpretation rules (fig. 2.4) and the duality requirement.

The Duality Requirement states that a constraint environment C is valid if there exist a substitution σ of variables ψ with closed session types, such that C_σ is well formed and for all $(c \sim \eta), (\bar{c} \sim \eta') \in C_\sigma$ we have $C \vdash \eta \bowtie \eta'$

The first algorithm is an adaptation of the homonymous algorithm from Amtoft, Neilsen & Neilsen [3]. From expression e if calculates it type, behaviour and constraint set. No session information is calculated.

The second algorithm infers a substitution and a refined set of constraints in such a way that the empty stack and behaviour τ will be reached when the rules from (fig. 2.4) are applied to a behaviour on which the substitution has been applied.

The third algorithm deals with the constraints relating to channels and generates duality constraints. It also checks these constraints.

Chapter 3

My Contribution

My contributions to this project are based on the second level. These include the development of a text based syntax for the intermediate code (the code produced from the first level), a lexer and parser for this syntax and a program to verify the behaviours produced.

The combination of lexer and parser converts the input code, which consists of one behaviour and one constraint, into the types defined in OCaml to represent the behaviours and constraints. A tuple of these is then be returned and passed into the behaviour checker where the rules from (fig. 2.4) are be applied to verify that the program communicates correctly.

One limitation of the checker that is worth noting is that it will not infer types and so these must be explicitly stated in the input code.

3.1 Text Based Syntax

3.1.1 Design

The syntax given in the paper can be seen in (fig. 2.1). To develop an easy to parse version of this syntax Greek letter have been removed, keywords or individual letters have been chosen to replace them. The final syntax is very similar to that given in (fig. 2.1). Substitutions for variables and labels are detailed in table 3.1. Keywords and the structure of the input can be found in the grammar detailed in 3.2.4.

Type Variables	α	[T][A-Z a-z 1-9]+
Behaviour Variables	β	[B][A-Z a-z 1-9]+
Session Variables	ψ	[S][A-Z a-z 1-9]+
Region Variables	ρ	[R][A-Z a-z 1-9]+
Labels	l	$[A-Z a-z 1-9]+$
channel	c	[C][A-Z a-z 1-9]+
channelEnd	$\bar{c}$	[T][A-Z a-z 1-9]+ [']

Table 3.1: Substitutions to given syntax

3.2 Lexer and Parser

Initially Camlp4 [5] was considered for this implementation. This is a Pre-Processor-Pretty-Printer for OCaml. However it became clear that this was not the appropriate tool since this is designed to extend the existing syntax of OCaml and not for developing new syntax structures.

3.2.1 Lexer

The lexer was designed using OCamlLex which is based on Lex, a lexical analyser generator. This takes a set of regular expressions and corresponding semantic actions. In this case the regular expressions are the keywords, labels, variables, etc. and the associated actions are either tokens that link to the parser or, in the case of the opening $, calls a second lexing function to deal with labels.

3.2.2 Parser

The parser was implemented using Menhir [6]. This is a parser generator that is closely related to OCamlyacc which in turn is based on yacc. These generate LR(1) parsers which means that the input is parsed from the bottom up and that there is one character look ahead.

To generate the parser the tokens are declared first. If they consist of a default type this is specified. For example the token $LABEL < string >$ states that LABEL consists of a string. The structure of the accepted input is then specified. With this the code that is to be invoked when this input is encountered is also stated. The structure of the accepted input can be found in 3.2.4.

3.2.3 Challenges

The main challenge encountered when designing the lexer and the parser was learning the correct usage of OCamlLex and Menhir. Having never used anything like this before from scratch these took some time to get used to. The tutorials [7][7][8][9, Chapter 16] were helpful in gaining an understanding of the syntax and structure.

Another challenge encountered with the lexer was parsing the end of file. An unknown issues causes the lexer to state that it finds a syntax error at the final character of the input file. While this is misleading to the user and not entirely correct it does not effect the behaviour of the behaviour checker and so has been ignored.

Finally initially there were some errors with shift/reduce conflicts in the parser. These were caused by not initially including operator precedence for ';' and ',' meaning that the execution order of sequencing of behaviours and constraints was unclear. As well as this there were not initially brackets in the

$$rec < behaviourVariable > (< behaviour >)$$

rule which also caused an issue. This was that it was unclear in the case

$$rec < behaviourVariable >< behaviour >; < behaviour >$$

if the recursive behaviour was in the sequence or if the sequence was the behaviour associated with the recursion.

3.2.4 Grammar

The grammar was designed with usability in mind. The main entry point for the grammar is shown first. This states that the parser will read a behaviour followed by a constraint followed by the end of file. In the description of the grammar all keywords, brackets etc. are input as they are given here. Anything inside <> is either another type from the grammar or a variable the syntax of which is given in table 3.1.

$\langle parse\_behaviour \rangle ::= \langle behaviour \rangle \langle constraint \rangle$ EOF

Behaviour

The behaviour type in the grammar consists of multiple options. The final option, for example, represents the offer of external choice and we can see that it contains a list. This will call to a sub grammar that states that this list consists of comma separated values of another subtype (opt_feild).

⟨*behaviour*⟩ ::= ⟨*behaviourVariable*⟩
| tau
| ⟨*behaviour*⟩ ; ⟨*behaviour*⟩
| chc (⟨*behaviour*⟩,⟨*behaviour*⟩)
| rec ⟨*behaviourVariable*⟩ (⟨*behaviour*⟩)
| spn (⟨*behaviour*⟩)
| psh (⟨*lable*⟩,⟨*sessionType*⟩)
| ⟨*region*⟩ ! ⟨*behaviourVariable*⟩
| ⟨*region*⟩ ? ⟨*behaviourVariable*⟩
| ⟨*region*⟩ ! ⟨*region*⟩
| ⟨*region*⟩ ? ⟨*lable*⟩
| ⟨*region*⟩ ! ⟨*lable*⟩
| ⟨*region*⟩ ? optn[⟨*oplist*⟩]

⟨*oplist*⟩ ::= (⟨*lable*⟩ ; ⟨*behavioiur*⟩) ⟨*opt_feild*⟩

⟨*opt_feild*⟩ ::= , (⟨*lable*⟩ ; ⟨*behavioiur*⟩)
| ϵ

Constraints

Constraints are similar to behaviours but also make use of the sub-grammar for bTypes, regions and session types which are given below.

⟨*constr*⟩ ::= ⟨*bType*⟩ <⟨*bType*⟩
| ⟨*behaviour*⟩ <⟨*behaviour*⟩
| ⟨*region*⟩ ∼ ⟨*regionVar*⟩
| ⟨*channel*⟩ ∼ ⟨*sessionType*⟩
| ⟨*channelEnd*⟩ ∼ ⟨*sessionType*⟩

| ⟨*contr*⟩ , ⟨*constr*⟩
| ϵ

Types

The implementation of types is relatively simple. They can be any of the forms listed below and use the sub-grammar for regions.

⟨*bType*⟩ ::= unit
| bool
| int
| pair (⟨*bType*⟩ ; ⟨*bType*⟩)
| funct ⟨*bType*⟩ ->⟨*bType*⟩ - ⟨*behaviourVariable*⟩
| ses ⟨*region*⟩
| ⟨*TVar*⟩

Region Variables

Regions are very simple and either consist of a region variable or a label.

⟨*regionVar*⟩ ::= ⟨*lable*⟩
| ⟨*region*⟩

Session Types

Session types follow a similar patter to behaviours.

⟨*sessionType*⟩ ::= end
| ! ⟨*bType*⟩ ⟨*sessionType*⟩
| ? ⟨*bType*⟩ ⟨*sessionType*⟩
| ! ⟨*sessionType*⟩ ⟨*sessionType*⟩
| ? ⟨*sessionType*⟩ ⟨*sessionType*⟩
| (+) [⟨*sesOpL*⟩] (⟨*lable*⟩ ; ⟨*sessionType*⟩)
| + [⟨*sesOpL*⟩] [⟨*sesOpL*⟩]
| ⟨*SVar*⟩

⟨*sesOpL*⟩ ::= (⟨*lable*⟩ ; ⟨*sessionType*⟩) ⟨*ses_opt_field*⟩

$\langle ses\_opt\_field \rangle$:= , ($\langle lable \rangle$; $\langle sessionType \rangle$)
| ϵ

3.3 OCaml Types

In order for the parser to have the correct types for dealing with behaviours and constraints these first needed to exist. This involved creating new types in OCaml for each of Behaviours, Constraints, Types, Regions and Session types. These main types are built up of subtypes, strings and, in the case of 'Tau' and 'End' nothing.

These types take the following form:

```
type b =
  | BVar of string
  | Tau
  | Seq of seq
  | ChoiceB of choiceB
  | RecB of recB
  | Spawn of spawn
  | Push of push
  | SndType of outT
  | RecType of recT
  | SndReg of sndR
  | RecLab of recL
  | SndChc of sndC
  | RecChoice of recC
  (* None included due to requirements to run main file *)
  | None
and sndC = { regCa : string ; labl : string}
and recC = { regCb : string ; cList : (string * b) list }
and recL = { regL : string ; label : string}
and sndR = { reg1 : string ; reg2 : string}
and recT = { regionR : string ; outTypeR : t}
and outT = { regionS : string ; outTypeS : t}
and push = { toPush : stackFrame}
and spawn = { spawned : b}
```

```
and recB = { behaVar : string ; behaviour : b}
and choiceB = {opt1 : b ; opt2 : b}
and seq = {b1 : b ; b2 : b};;
```

You can clearly see how RecChoice, the external choice type is made up of the subtype recC which in turn consists of a string and list of string, behaviour tuples. The other types are implemented in a similar way.

As well as implementing the types each type also has a to_string function. This is to allow for the re-printing of the input code as part of the output.

3.3.1 Challenges

Again the main challenges with implementing this part of the project was the new language. While OCaml has excellent documentation it was not easy to find examples or tutorials to help with implementing such an interlinked system of types. One tutorial that was helpful was [10].

3.4 Behaviour Checker

The behaviour checker is implementing the rules from (fig. 2.4). The first step towards implementing these was to store the relevant information in an easily accessible form. In this case that meant storing the relevant constraints. Then the behaviours have to be checked in relation to these constraints and the relevant actions from the rules applied.

3.4.1 Storing the constraints

The constraints that are relevant to the Behaviour Checker are the region constraints, the behaviour constraints and the type constraints. Each is stored in a slightly different way to account for the fact that the format and usage of each of the constraint types is different. Other constraints are not stored.

Behaviour constraints

These constraints are of the form

$$b \subseteq \beta$$

where b is a behaviour and β is a behaviour variable. These constraints are used in the Beta and Rec rules. In the Beta rule ($\Delta \models \beta \rightarrow c\Delta \models b$ if $C \vdash b \subseteq \beta$) the constraints are used to replace the behaviour variable beta with each behaviour associated with it. Each of these is then checked against the rules to see if it produces a valid end state consisting of the empty stack and the τ behaviour.

In the recursion (Rec) rule the behaviour constraints are used where there are any constraints of which the left hand side matches the current recursion behaviour. If any are found then the left hand side of that rule is replaced with τ.

The decision was made to store the behaviour constraints as a hash table. The key is the right hand side of the constraint β and the value is the left hand side (the behaviour associated with β). In this way in the case of either rule we can search quickly and find all values associated with the key and have them returned as a list. We can then either replace β with each value from the list in turn (Beta Rule) or we can search the list to find and replace any instances of the current recursion behaviour (Rec Rule).

Type constraints

Type constraints are of the form

$$T_1 <: T_2$$

and are used in the Out Rule as well as in endpoint relation checks in the Delegate Rule. They are used according to semantics given in (fig. 3.1) that state that if the constraint exists then T_1 is a functional subtype of T_2. If both T_1 and T_2 are pairs then if the first type of the first pair is a subtype of the first type of the second pair and if the same is true of the second types of both pairs then the first pair is a subtype of the second. It is a similar case for function.

If both types are of the form Ses^{ρ} then if there exists a region constraint linking the ρ's they are functional subtypes.

Type constraints are stored as a hash table with the right hand type as the key and

$$\frac{(T_1 \subseteq T_2) \in C}{C \vdash T_1 <: T_2} \quad \frac{C \vdash \rho \sim \rho'}{C \vdash Ses^{\rho} <: Ses^{\rho'}} \quad \frac{C \vdash T_1' <: T_1; C \vdash \beta \subseteq \beta; C \vdash T_2 <: T_2'}{C \vdash T_1 \xrightarrow{\beta} T_2 <: T_1' \xrightarrow{\beta'} T_2'}$$

$$\frac{}{C \vdash T <: T} \quad \frac{C \vdash T_1 <: T_2; C \vdash T_2 <: T_3}{C \vdash T_1 <: T_3} \quad \frac{C \vdash T_1 <: T_2; C \vdash T_3 <: T_4}{C \vdash T_1 \times T_3 <: T_2 \times T_4}$$

Figure 3.1: Functional Subtyping

the left hand type as the value. Since type constraints are also transitive this means that when checking these constraints first the right hand side of the constraint to be checked is searched for. A list of associated subtypes for this type are returned. Each of these is then checked to see if it matches the left hand side. If not the search is repeated with this new type as the right hand side. This continues until either the left hand side is found or there are no more types to check.

Since type constraints are also reflexive it is always checked if the two types are equal first.

Region constraints

Region constraints take the form of either

$$\rho \sim \rho' \; or \; \rho \sim l$$

and are used in the rules: Out, In, Del, Res, ICh and ECh. They are used in the same way for each of the rules which is simply to check if there is a link between a region and a label. The complication arises from the fact that regions can be chained and then linked to a label. For example $\rho_1 \sim \rho_2$, $\rho_2 \sim \rho_3$ and $\rho_3 \sim l1$ tells us that ρ_1, ρ_2 and ρ_3 are all linked to $l1$.

Due to this region constraints are stored as a list of tuples where each tuple consist of a label and a hash table with region variables as keys and unit as the value. In this way when looking up a particular label and region we can simply search for the label in the list and then check the hash table. This is an efficient implementation since the number of labels in any given program is not likely to be excessive.

Storing the constraints in this way in a single pass through the constraint list was challenging. It is achieved by first creating a new tuple for the next constraint to be stored. This consists of either a label or a None and a hash table with either one or two values depending on the structure of the constraint. All other elements from the list that match either of the two values associated with the new constraint are then found and removed from the list. They are merge together with the new element to form a single list element which is then added back to the list.

3.4.2 Behaviour checker function

The behaviour checker verifies that the input program will in fact communicate correctly. It takes an input of the behaviours produced from the first level (2.2) and the constraints stored in the method described above. The behaviours are then checked against the rules from (fig. 2.4) and the appropriate actions taken.

The parameters to the function are a set of constraints and a list of behaviour tuples of the form $(behaviour, stack, stack\_labels, continuation)$. The stack represents the one described in the rules and the stack labels structure is used to ensure that each label is only ever pushed to the stack once. The continuation is used to keep track of any behaviours that need to be dealt with once the current behaviour is finished with, see (fig. 3.2) for an example.

The initial function called is a wrapper function that calls to a check step function with the first tuple from the list.The check_step function then takes the current stack and checks to see if the top frame contains the 'End' session type (i.e. it applies the end rule) if it does it removes this frame and then it continues to check the to see which of the rest of the rules apply. The checks are performed in this order since the End Rule is performed based on the state of the stack while all other rules are performed based on the current behaviour and the state of the stack.

3.4.3 Results

The behaviour checker will output for the results of the check. In the case of the check failing it will also print the name of the rule on which the check failed and if it failed due to a failed constraint check. It will not print the location of the error in the source code.

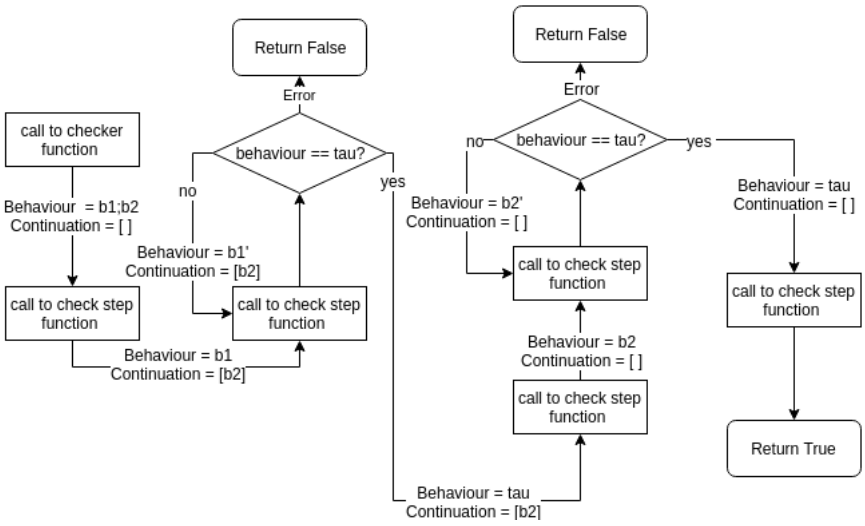

Figure 3.2: Application of Sequence Rule

If the check is successful we have shown that the input code follows the correct protocol for the communication channel. We have also shown partial lock freedom due to the well stackness of the program.

3.4.4 Challenges

The main challenges encountered in this implementation of the checker was, again, the new language. However this was overcome more quickly when developing this section of the project.

Other challenges encountered included dealing with the different implementations of stack and hash tables in the standard library and the core_kernal library. The stacks in the core_kernel library included functions that made for a nicer implementation but the hash tables were missing the functionality to return lists of all bindings to a key. This was overcome when I realised it was possible to rename the hash table module from the standard library before importing the core_kernel library and so avoid the shadowing of the binding.

The implementation of the region constraint storage also took quiet some time since it is rather complicated code and pushed my knowledge of the OCaml language to its boundaries.

Chapter 4

Motivating Examples

In this chapter some sample programs are shown and the process through which they are run is explained.

4.1 Simple Swap Service

The examples given in the paper [1] include an example for a simple swap service. This consists of two clients and a coordinator. The clients send values to the coordinator who then returns to client one the value client two and vice versa.

4.1.1 Input code

```
let fun coord(_) =
  let val p1 = acc swp ()
      val x1 = recv p1
      val p2 = acc swp ()
      val x2 = recv p2
  in send p2 x1; send p1 x2;
     coord ()
in spawn coord;
```

```
let fun swap(x) =
  let val p = req swp ()
  in send p x; recv p
in spawn (fn _ => swap 1);
   spawn (fn _ => swap 2);
```

In this code the coordinator accepts two connects on the channel swp using acc-swp (). This gives it two endpoints, p1 and p2. After accepting each connections it receives

a value over that connection. Finally the received values are sent over the alternate endpoint.

In the swap processes a connection is requested on the swp channel using req-swp (), giving the endpoint p. A value, in this case either 1 or 2, is then sent over this endpoint and finally some other value is received over p.

The endpoints used by the coordinator, received from calling acc-swp (), are used according the the session type $?T.!T.end$. This means that on this connection a value of type T is received, then a value of the same type is sent then the connection ends.

The endpoint used by the swap function, received from the call to req-swp (), is used according to the session type $!Int.?T'.end$. In this case an int is received, a value of some type T' is sent and then the connection is terminated.

From examining these session types and the above code it is clear that types T and T' must be the same. This is due to the fact that the values sent by the coordinator are then received by the swap functions. Furthermore both of these types must be of the type *int* since the values received by the coordinator are sent by the swap functions and these values are integers (1 and 2).

In this case, the programs can communicate and are typable when $T = T' = Int$. The inference algorithms proposed in (Spaccasassi & Koutavas)[1] can infer these session types from the code given.

4.1.2 Intermediate code

The first level of the system will produce code detailing the behaviour of the given input code and the constraints relevant to it. A simplified version of this can be seen here 4.1 and the full version in the appendix.

This code is given in the syntax developed by me for this project (see 3.1). In this case the first spawn of behaviour variable B101 is referencing the spawning of the coordinator in the input code. Behaviour variable B101 then references the behaviour of the coordinator. The second and third spawns reference the spawns of the swaps.

In the constraints we can first see the three channel constraints. These are referring to the endpoints returned from the calls to acc-swp and req-swp. The next constraint is a behaviour constraint. This is what gives us the link between the behaviour variable B101 and the actual behaviour of the coordinator. Finally the last constraints are the

Behaviours:

```
spn (B101);
spn (psh (l3, S126);R131 ! int;R132 ? T125);
spn (psh (l3, S126);R147 ! int;R148 ? T141)
```

Constraints:

```
Cswap1' ~ S105,
Cswap2' ~ S106,
Cswap3 ~ S126,
rec B101(psh (l1, S105);R114 ? T103 ;psh (l2, S106);R115 ? T104;R116
    ! T103;R117 ! T104;B101) < B101,
l1 ~ R114,
l1 ~ R117,
l2 ~ R115,
l2 ~ R116,
l3 ~ R131,
l3 ~ R132,
l3 ~ R147,
l3 ~ R148
```

Figure 4.1: Simplified intermediate code for Simple Swap

region constraints. These either link regions to other regions or labels to regions. The labels here are static approximations of the locations of the endpoints in the code. In this case $l1$ refers to the location in the coordinator function where acc-swp is called to generated the connection for endpoint p1. $l2$ references the location of p2 and $l3$ references the call to req-swp in the spawn function.

Since the behaviour checker does not deal with inference we must manually apply some substitutions before we can run this code through the it. This involves replacing the session variables with the actual session endpoint types and the type variables with the types. These substitutions are detailed in table 4.1.

4.1.3 Behaviour check

The intermediate code with substitutions can then be run through the behaviour checker detailed in 3.4. Here the steps which would be taken to check the simplified code are outlined.

First the code is run through the parser and lexical analyser. The resulting con-

S105 S106	?int !int end
S126	!int ?int end
T125 T141 T103 T104	int

Table 4.1: Manual Substitutions for Simple Swap

straints are then passed into the function that deals with their internal storage. The simplified output from running this example, which includes the printing of the internal storage of the constraints, can be seen in 4.2. It can clearly be seen here how the internal storage allows for efficient look ups, particularly with respect to region constraints. The stored constraints and the behaviours are then passed to the behaviour checker.

The first behaviour is a sequence. Once the rule is applied to deal with this the next behaviour to look at is the $spn(B101)$. At this stage the stack is empty. This behaviour matches the rule SPAWN from (fig. 2.4). This then calls the check function again with the behaviour tau and the result is anded with the result of a call to the check functions where the behaviour is the body of the spawn ($B101$).

The rules then continue to be applied in such a fashion until a state is reached where the stack is empty, the continuation is empty and the behaviour is tau. In this case the check function returns true. Alternately if a state is reached where no rule can be applied false is returned along with a message detailing where in the rules the error state was reached. A detailed diagram of how the simplified code would be dealt with in the behaviour checker can be found in (fig. 4.3).

```
Behaviours:
Spn( B101 );
Spn( Psh( l3, ! int ? int end ); R131 ! int; R132 ? int);
Spn( Psh( l3, ! int ? int end ); R147 ! int; R148 ? int)

Paired:
  Beta: B101
  Behaviour: rec B101 (Psh( l1, ? int ! int end ); R114 ? int; Psh
     ( l2, ? int ! int end ); R115 ? int; R116 ! int; R117 ! int;
     B101)

Region Constraints:
 label: l3
 regions:
  R132, R148, R131, R147
label: l2
 regions:
  R116, R115
label: l1
 regions:
  R117, R114

check successful!
```

Figure 4.2: Output from Behaviour Checker for Simple Swap

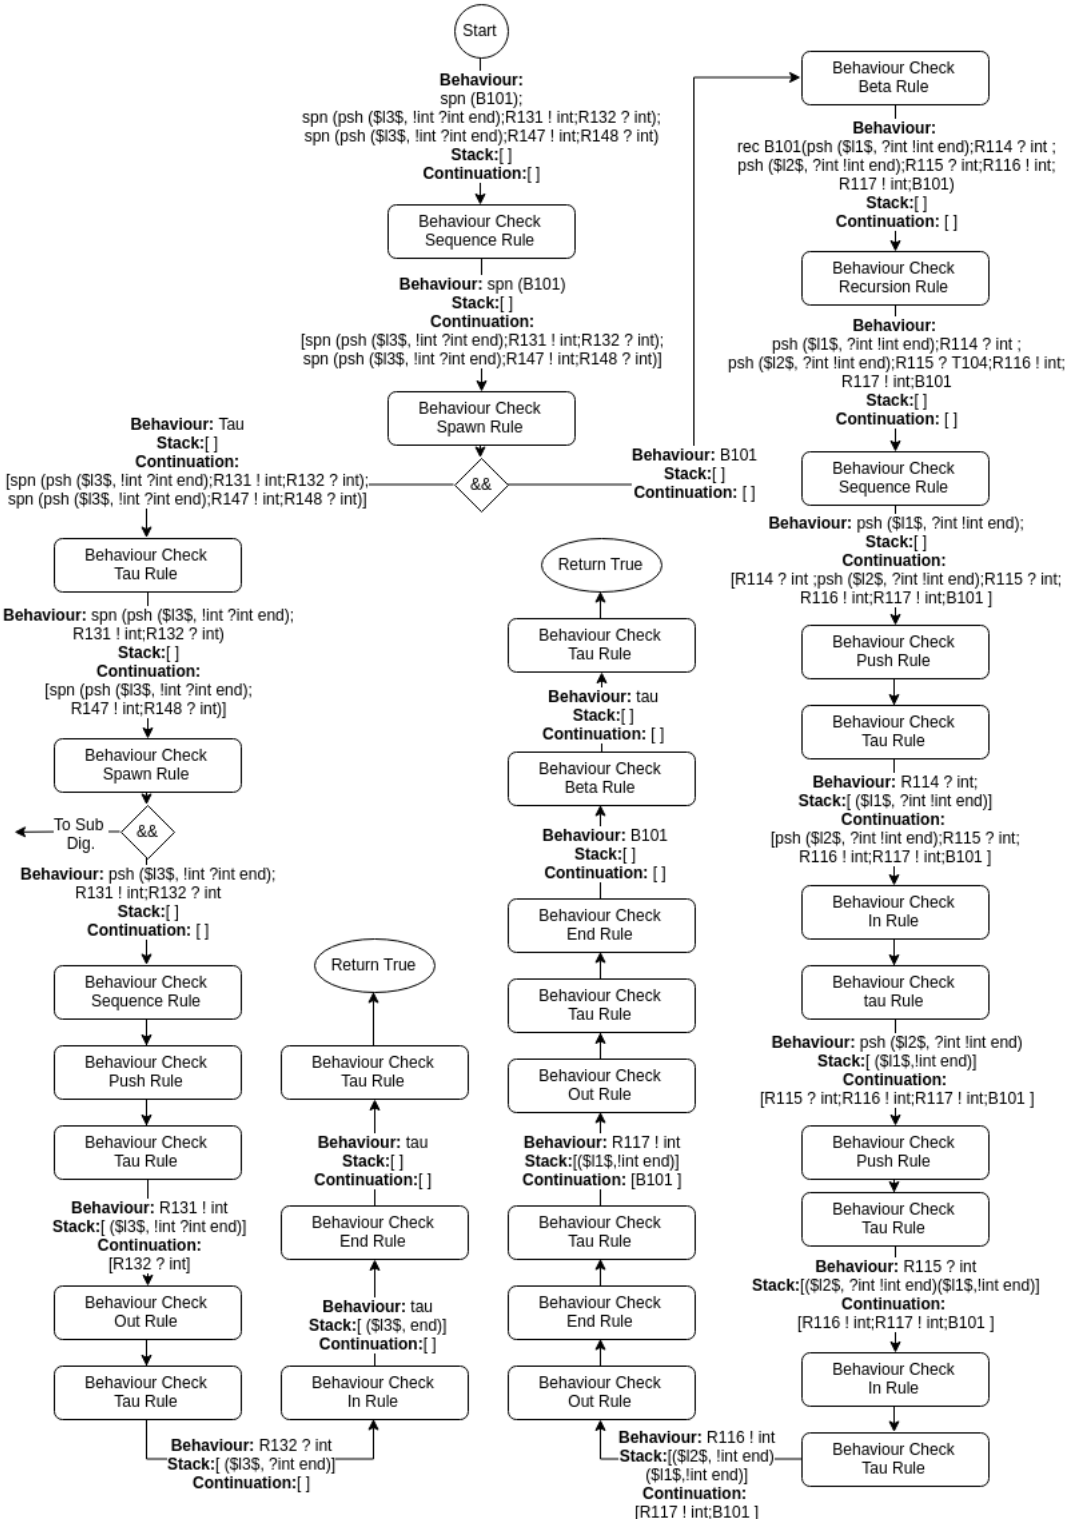

Figure 4.3: Simple Swap Path Taken Through Behaviour Check, Main Dig.

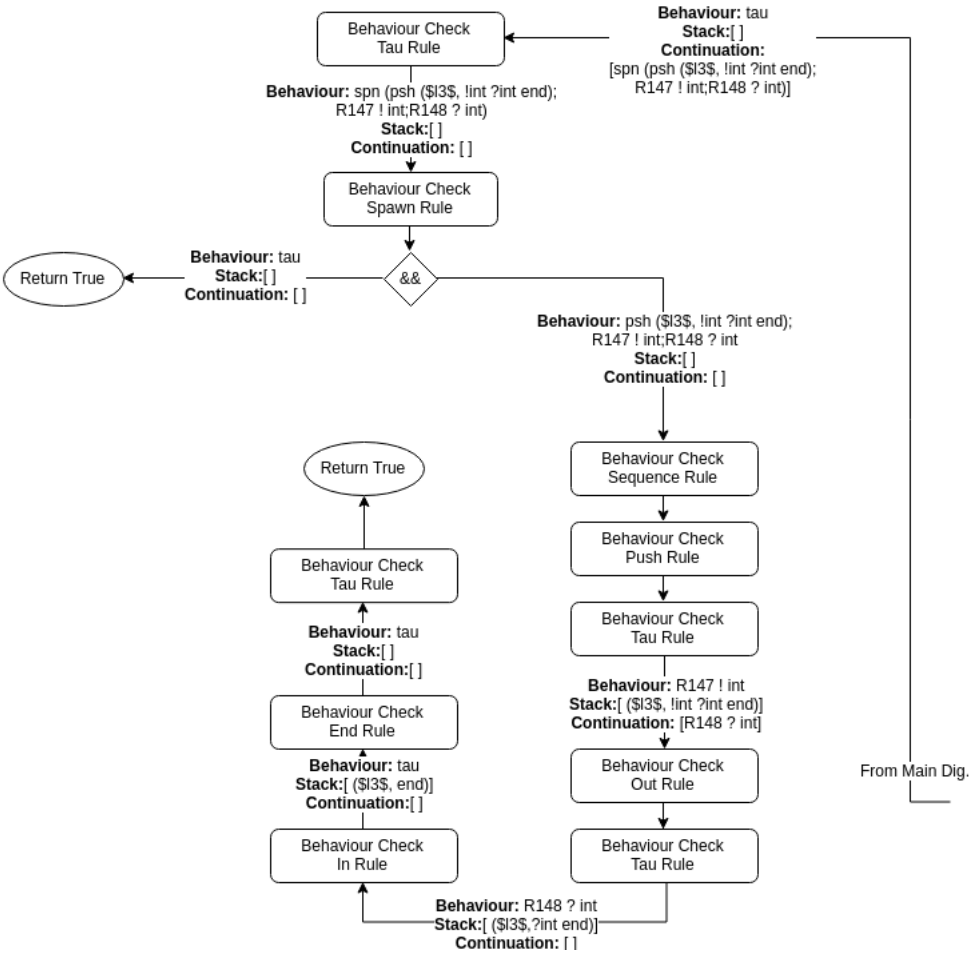

Figure 4.4: Simple Swap Path Taken Through Behaviour Check, Sub Dig.

4.2 Swap Delegation

This example is given in the paper (Spaccasassi & Koutavas)[1]. Here it has been altered and extended to show how incorrect code would be dealt with by the behaviour checker.

In this example the coordinator function delegates the exchange to the clients. In the previous example the coordinator function is a bottleneck when there are large exchange values, the delegation in this example avoids this.

4.2.1 Input code

Here the swap function connects to the coordinator using req-swp as before but now it offers two choices; SWAP and LEAD. If SWAP is selected by the coordinator then the function receives on the endpoint p then sends the value x over p.

Otherwise if LEAD is selected the function resumes the endpoint q. This endpoint is received over p. A value is the received on q and the value passed in to the function is then sent over q.

The coordinator accepts two connections on the swp channel and gets the two endpoints p1 and p2. It then selects SWAP on p1 and LEAD on p2. Then p1 is sent over p2 and the function recurs.

```
let fun coord(_) =
  let val p1 = acc swp ()
  in sel SWAP p1;
    let val p2 = acc swp
    in sel LEAD p2;
       deleg p2 p1;
       coord()
```

```
let fun swap(x) =
  let val p = req swp ()
  in case p {
    SWAP: recv p; send p x
    LEAD: let val q = resume
        p
              val y = recv q
          in send q x; y }
```

When this coordinator function is analysed in isolation the inference algorithms will infer that the endpoints p1 and p2 have the type $\eta_{i,i\in(p1,p2)} = (\oplus\{(SWAP : \eta'), (LEAD : \eta'.end)\})$. When swap is analysed they will infer that $\eta_p = \Sigma\{(SWAP : ?T_1.!T_2.end), (LEAD :?\eta_q.end)\}$ and $\eta_q =?T_2.!T_1.end$.

When the coordinator is looked at in isolation then η' can be any endpoint. However due to duality η' must be equal to η_q in this case. Also $T_1 = T_2$ must be true.

S120 S121	$(+)[(\text{SWAP}; ?int!intend), (\text{LEAD}; !?int!intendend)]$
S142	$+[(\text{SWAP}; ?int?intend), (\text{LEAD}; ??int!intendend)][\,]$
T160 T191 T157 T188	int

Table 4.2: Manual Substitutions for Delegated Swap

The deliberate error in this code is that in the swap function when SWAP is selected a value is received then a value is sent. When LEAD is selected the same pattern is followed. This will produce a deadlock where both instances of swap are waiting on the other to send a value.

4.2.2 Intermediate code

A simplified version of the output from the first level of the system is given in (fig. 4.5).

Again substitutions must be made before this code can be given to the behaviour checker. These are detailed in table 4.2.

4.2.3 Behaviour check

As you can see from (fig. 4.6) the behaviour checker has recognised that the input code cannot communicate correctly. The error given is ‘

$$outrulestackframeincorrectforcurrentbehaviour$$

’. This tells us that the checker ran into problems when it was trying to verify a behaviour of the type $\rho!T$. We also know, from looking up the rules, that the stack frame did not match the expected form $(l; !T\eta)$. From these we can deduce that the checker ran into the error when attempting to check if (fill in)

Behaviours:

```
spn (B118);
spn (psh (l3, S142);R150? optn [(SWAP ; R152 ! int;R151 ? T160),
   (LEAD ; R153 ? l4;R154 ? T157;R155 ! int)]);
spn (psh (l3, S142);R181? optn [(SWAP ; R183 ! int;R182 ? T191),
   (LEAD ; R184 ? l4 ;R185 ? T188;R186 ! int)])
```

Constraints:

```
unit <  ses R151,
unit <  ses R182,
Cswap1' ~ S120,
Cswap2' ~ S121,
Cswap3 ~ S142,
rec B118(psh (l1, S120);R128 ! SWAP;psh (l2, S121);R129 ! LEAD;
   R130 ! R131;B118) < B118,
l1 ~ R128,
l1 ~ R131,
l2 ~ R129,
l2 ~ R130,
l3 ~ R150,
l3 ~ R152,
l3 ~ R153,
l3 ~ R181,
l3 ~ R183,
l3 ~ R184,
l4 ~ R154,
l4 ~ R155,
l4 ~ R185,
l4 ~ R186
```

Figure 4.5: Simplified intermediate code for Delegated Swap

```
Behaviours:
   Spn( B118 );
   Spn( Psh( l3, +[ (SWAP;? int ? int end ), (LEAD;? ? int ! int
      end end ) ][]) ; R150 ? optn [(SWAP; R152 ! int ; R151 ? int
      ) (LEAD;  R153 ? l4 ;   R154 ? int ; R155 ! int) ] );
   Spn( Psh( l3, +[ (SWAP;? int ? int end ), (LEAD;? ? int ! int
      end end ) ][]  ); R181 ? optn [(SWAP; R183 ! int; R182 ? int
      ) (LEAD;  R184 ? l4 ; R185 ? int; R186 ! int) ] )
Behaviour constraints:
Beta: B118
Behaviour: rec B118 (Psh(l1, (+)[ (SWAP;? int ! int end ), (LEAD;!
    ? int ! int end end )]) ; R128 ! SWAP; Psh(l2, (+)[ (SWAP;?
   int ! int end ), (LEAD;! ? int ! int end end )]); R129 ! LEAD;
   R130 ! R131;B118)

Region Constraints:
 label: l4
 regions:
R154
R185
R186
R155
label: l3
 regions:
R152
R183
R184
R181
R150
R153
label: l2
 regions:
R129
R130
label: l1
 regions:
R128
R131

Type Constraints:
Paired:
Super: ses R182
sub:  unit
Paired:
Super: ses R151
sub:  unit

out rule
stack frame incorrect for current behaviour
ERROR: Failed Check
```

Figure 4.6: Output from Behaviour Checker for Delegated Swap

4.3 TSL

A very simple implementation of the Transport Layer Security (TSL) handshake has been detailed here to show how it would operate under this system. Again the simplified version of the output from the two layers has been shown here and the full version has been included in the appendix.

4.3.1 Input code

4.3.2 Intermediate code

4.3.3 Behaviour check

Chapter 5

Evaluation

This chapter explains how the application was evaluated as well as giving a reflection on the results of the evaluation.

5.1 Testing

Testing was conducted in stages. Firstly the parser and lexer were tested, then the constraint storage and finally the behaviour checker.

5.1.1 Lexer and parser

The testing for the lexer and parser was relatively simple. To string methods were written for each of the data types and the resulting behaviours and constraints from the parsing were printed out to the console.

Initially they were printed in the exact form that they were input. This allowed a side by side comparison of the input and output which allowed for errors to be spotted quickly and easily.

5.1.2 Constraint storage

Testing for the constraint storage was done in a similar way to that of the lexer and parser. Instead of printing the constraints as they were read in the function was

updated to print them as they were been stored. This allowed for them to be quickly checked against the input to see if they were been stored correctly.

For example in the case of region constraints the input could contain constraints such as:

```
R12 ~ l1,
R11 ~ R12,
R15 ~ l3
```

Which should then be output in the form:

```
label: l1
regions: R11, R12
label: l2
regions: R15
```

It is then easy to check if the regions listed under a particular label are in fact linked to it.

5.1.3 Behaviour checker

The testing for the behaviour checker was the most in depth. A test suit has been developed that includes the examples given in 4. As well as these several small programs were written to test the individual rules one at a time.

These small programs were first written to test each of the rules with the basic situations under which they should pass or fail. The scenarios for passing involve setting up a situation where a simple form of the correct stack frame and behaviour are examined. It is also ensured that the relevant constraints are met.

The test to ensure that the checker fails correctly involve setting up situations where each of the constraints for a rule is broken in turn. Ideally only one constraint is broken per test but in some cases it is more practical to test breaking multiple constraints at once.

Writing the test in this way has also shown that the constraint checks work as expected.

5.2 Reflection

The test detailed in the previous section have shown that all components of the program behave as is expected. While it is possible that more extensive test might find some unexpected behaviour it seems unlikely.

The development of the project in a functional style meant that once the files compiled correctly the behaviour of the program was almost always what was specified. In the main the problems discovered by the tests were my own misunderstandings of the rules or the constraints. The tests these misunderstandings more obvious and so easier to fix.

The test was also invaluable in showing where the error hints from the behaviour checker should display. For example the tests for the ICh rule showed an uncaught Not_found exception instead of displaying an error message. This allowed me to go back to the program, catch the exception and output a more helpful message.

Chapter 6

Conclusion

Throughout the duration of this project

6.1 Evaluation of Objectives

The objectives of this project were met. That is to say that an understanding of session types and the current attempt to verify them was achieved. In the initial plan for this project it was thought that time might allow for the implementation of the inference algorithms and their integration with the behaviour checker for a more complete implementation of the second level of the system from (Spaccassi & Koutavas)[1]. This however proved to be impractical in the time allocated.

The testing detailed in 5 shows that the implementation of the behaviour checker is complete and that it functions in the expected fashion.

6.2 Reflection on Process

The process of development of this project could perhaps have been improved. Starting out by spending a bit more time studying the paper and perhaps first writing some simple programs in OCaml might have saved some time on developing the behaviour checker. It would certainly have helped when technical difficulties with the language were encountered.

Appendix

Simple Swap

Full intmediate code

The substitutions to run this code through the behaviour checker are given in 4.1.

```
tau;tau;tau;spn (B101);tau;spn (B118);tau;spn (B134)

Cswap1' ~ S105,
Cswap2' ~ S106,
Cswap3 ~ S126,
tau < B41,
tau < B57,
tau < B85,
psh (l1, S105) < B108,
psh (l2, S106) < B110,
psh (l3, S126) < B128,
psh (l3, S126) < B144,
R116 ! T103 < B112,
R117 ! T104 < B113,
R131 ! int < B129,
R147 ! int < B145,
R114 ? T103 < B109,
R115 ? T104 < B111,
R132 ? T125 < B130,
R148 ? T141 < B146,
tau;tau;B123;tau < B118,
```

```
tau;tau;B128;tau;tau;tau;B129;tau;tau;B130 < B123,
tau;tau;B139;tau < B134,
tau;tau;B144;tau;tau;tau;B145;tau;tau;B146 < B139,
rec B101(tau;tau;B108;tau;tau;B109;tau;tau;B110;tau;tau;B111;tau;
   tau;tau;B112;tau;tau;tau;B113;tau;tau;B101) < B101,
l1 ~ R114,
l1 ~ R117,
l2 ~ R115,
l2 ~ R116,
l3 ~ R131,
l3 ~ R132,
l3 ~ R147,
l3 ~ R148
```

Full output from behaviour checker

Out from behaviour checker where substitutions are made prior to running.

```
Behaviours:
 Tau; Tau; Tau; Spn( B101 ); Tau; Spn( B118 ); Tau; Spn( B134 )

Behaviour constraints:
Paired:
Beta: B144
Behaviour: Psh ( l3, ! int ? int end )
Paired:
Beta: B113
Behaviour: R117 ! int
Paired:
Beta: B128
Behaviour: Psh ( l3, ! int ? int end )
Paired:
Beta: B146
Behaviour: R148 ? int
Paired:
Beta: B41
```

```
Behaviour: Tau
Paired:
Beta: B123
Behaviour: Tau; Tau; B128 ; Tau; Tau; Tau; B129 ; Tau; Tau; B130
Paired:
Beta: B57
Behaviour: Tau
Paired:
Beta: B101
Behaviour: rec B101 Tau; Tau; B108 ; Tau; Tau; B109 ; Tau; Tau;
   B110 ; Tau; Tau; B111 ; Tau; Tau; Tau; B112 ; Tau; Tau; Tau;
   B113 ; Tau; Tau; B101
Paired:
Beta: B118
Behaviour: Tau; Tau; B123 ; Tau
Paired:
Beta: B134
Behaviour: Tau; Tau; B139 ; Tau
Paired:
Beta: B139
Behaviour: Tau; Tau; B144 ; Tau; Tau; Tau; B145 ; Tau; Tau; B146
Paired:
Beta: B111
Behaviour: R115 ? int
Paired:
Beta: B110
Behaviour: Psh ( l2 , ? int ! int end )
Paired:
Beta: B129
Behaviour: R131 ! int
Paired:
Beta: B109
Behaviour: R114 ? int
Paired:
Beta: B85
```

```
Behaviour: Tau
Paired:
Beta: B108
Behaviour: Psh ( l1, ? int ! int end )
Paired:
Beta: B112
Behaviour: R116 ! int
Paired:
Beta: B145
Behaviour: R147 ! int
Paired:
Beta: B130
Behaviour: R132 ? int
Region Constraints:
 label: l3
 regions:
  R132
  R148
  R131
  R147
label: l2
 regions:
  R116
  R115
label: l1
 regions:
  R117
  R114

Type Constraints:

check successful!
```

Swap Delegation

Full intermeditate code

The substitutions for this code to be run through the behaviour checker are given in 4.2.

```
tau;tau;tau;spn (B118);tau;spn (B132);tau;spn (B163)

unit <  ses R151,
unit <  ses R182,
Cswap1' ~ S120,
Cswap2' ~ S121,
Cswap3 ~ S142,
tau < B38,
tau < B74,
tau < B102,
psh (l1, S120) < B123,
psh (l2, S121) < B125,
psh (l3, S142) < B144,
psh (l3, S142) < B175,
R152 ! int < B146,
R155 ! int < B149,
R183 ! int < B177,
R186 ! int < B180,
R151 ? T160 < B145,
R154 ? T157 < B148,
R182 ? T191 < B176,
R185 ? T188 < B179,
R130 ! R131 < B127,
R153 ? l4 < B147,
R184 ? l4 < B178,
R128 ! SWAP < B124,
R129 ! LEAD < B126,
tau;tau;B137;tau < B132,
```

tau;tau;B144;tau;R150? optn {($SWAP$; tau;tau;tau;tau;tau;B146; B145), ($LEAD$; tau;tau;B147;tau;tau;B148;tau;tau;tau;B149;tau) } < B137,
tau;tau;B168;tau < B163,
tau;tau;B175;tau;R181? optn {($SWAP$; tau;tau;tau;tau;tau;B177; B176), ($LEAD$; tau;tau;B178;tau;tau;B179;tau;tau;tau;B180;tau) } < B168,
rec B118(tau;tau;B123;tau;B124;tau;tau;B125;tau;B126;tau;tau;tau; B127;tau;tau;B118) < B118,
$l1$ ~ R128,
$l1$ ~ R131,
$l2$ ~ R129,
$l2$ ~ R130,
$l3$ ~ R150,
$l3$ ~ R152,
$l3$ ~ R153,
$l3$ ~ R181,
$l3$ ~ R183,
$l3$ ~ R184,
$l4$ ~ R154,
$l4$ ~ R155,
$l4$ ~ R185,
$l4$ ~ R186

Full output from behaviour checker

```
Behaviours:
Tau; Tau; Tau; Spn( B118 ) ; Tau; Spn( B132 ) ; Tau; Spn( B163 )
Behaviour constraints:

Paired:
Beta: B132
Behaviour: Tau; Tau; B137 ; Tau
Paired:
```

```
Beta: B102
Behaviour: Tau
Paired:
Beta: B144
Behaviour: Psh ( l3 , +[ (SWAP;? int ? int end )  (LEAD;? ? int !
   int end end ) ][ ]  )
Paired:
Beta: B137
Behaviour: Tau; Tau; B144 ; Tau; R150 ? optn [(SWAP;Tau; Tau; Tau;
    Tau; Tau; B146 ; B145) (LEAD;Tau; Tau; B147 ; Tau; Tau; B148 ;
    Tau; Tau; Tau; B149 ; Tau) ]
Paired:
Beta: B38
Behaviour: Tau
Paired:
Beta: B179
Behaviour: R185 ? int
Paired:
Beta: B126
Behaviour: R129 ! LEAD
Paired:
Beta: B125
Behaviour: Psh ( l2 , (+)[ (SWAP;? int ! int end )  (LEAD;! ? int !
    int end end ) ]  )
Paired:
Beta: B146
Behaviour: R152 ! int
Paired:
Beta: B123
Behaviour: Psh ( l1 , (+)[ (SWAP;? int ! int end )  (LEAD;! ? int !
    int end end ) ]  )
Paired:
Beta: B177
Behaviour: R183 ! int
Paired:
```

```
Beta: B124
Behaviour: R128 ! SWAP
Paired:
Beta: B168
Behaviour: Tau; Tau; B175 ; Tau; R181 ? optn [(SWAP;Tau; Tau; Tau;
    Tau; Tau; B177 ; B176) (LEAD;Tau; Tau; B178 ; Tau; Tau; B179 ;
    Tau; Tau; Tau; B180 ; Tau) ]
Paired:
Beta: B176
Behaviour: R182 ? int
Paired:
Beta: B127
Behaviour: R130 ! R131
Paired:
Beta: B118
Behaviour: rec B118 Tau; Tau; B123 ; Tau; B124 ; Tau; Tau; B125 ;
   Tau; B126 ; Tau; Tau; Tau; B127 ; Tau; Tau; B118
Paired:
Beta: B74
Behaviour: Tau
Paired:
Beta: B149
Behaviour: R155 ! int
Paired:
Beta: B180
Behaviour: R186 ! int
Paired:
Beta: B147
Behaviour: R153 ? l4
Paired:
Beta: B178
Behaviour: R184 ? l4
Paired:
Beta: B163
Behaviour: Tau; Tau; B168 ; Tau
```

```
Paired:
Beta: B148
Behaviour: R154 ? int
Paired:
Beta: B175
Behaviour: Psh ( l3 , +[ (SWAP;? int ? int end )  (LEAD;? ? int !
   int end end ) ][ ]  )
Paired:
Beta: B145
Behaviour: R151 ? int
Region Constraints:
 label: l4
 regions:
R154
R185
R186
R155
label: l3
 regions:
R152
R183
R184
R181
R150
R153
label: l2
 regions:
R129
R130
label: l1
 regions:
R128
R131

Type Constraints:
```

```
Paired:
Super: ses R182 sub:
 unit
Paired:
Super: ses R151 sub:
 unit
out rule
stack frame incorrect for current behaviour
ERROR: Failed Check
```

Bibliography

[1] C. Spaccasassi and V. Koutavas, "Type based analysis for session inference."

[2] H. Hüttel, I. Lanese, V. T. Vasconcelos, L. Caires, M. Carbone, P.-M. Denilou, D. Mostrous, L. Padovani, A. Ravara, E. Tuosto, H. T. Vieira, and G. Zavattaro, "Foundations of behavioural types." `http://www.behavioural-types.eu/publications`, 2014.

[3] T. Amtoft, F. Nielson, and H. R. Nielson, *Type and Effect System - Behaviours for Concurrency*. Imperial College Press, 1999.

[4] G. Castagna, M. Dezani-Ciancaglini, E. Giachino, and L. Padovani, "Foundations of session types," 2009.

[5] OCaml. `https://github.com/ocaml/camlp4/wiki`.

[6] F. Pottier and Y. Régis-Gianas, "Menhir reference manual." `http://pauillac.inria.fr/~fpottier/menhir/manual.pdf`.

[7] T. Team. `http://toss.sourceforge.net/ocaml.html`, 2013.

[8] Yan. `http://yansnotes.blogspot.ie/2014/11/menhir.html`, 2014.

[9] Y. Minsky, A. Madhavapeddy, and J. Hickey, *Real World OCaml*. O'Reilly Media, 2013.

[10] *Functional programming using Caml Light*, ch. 6.